


Disaster Tweets on A10: two precision strategies

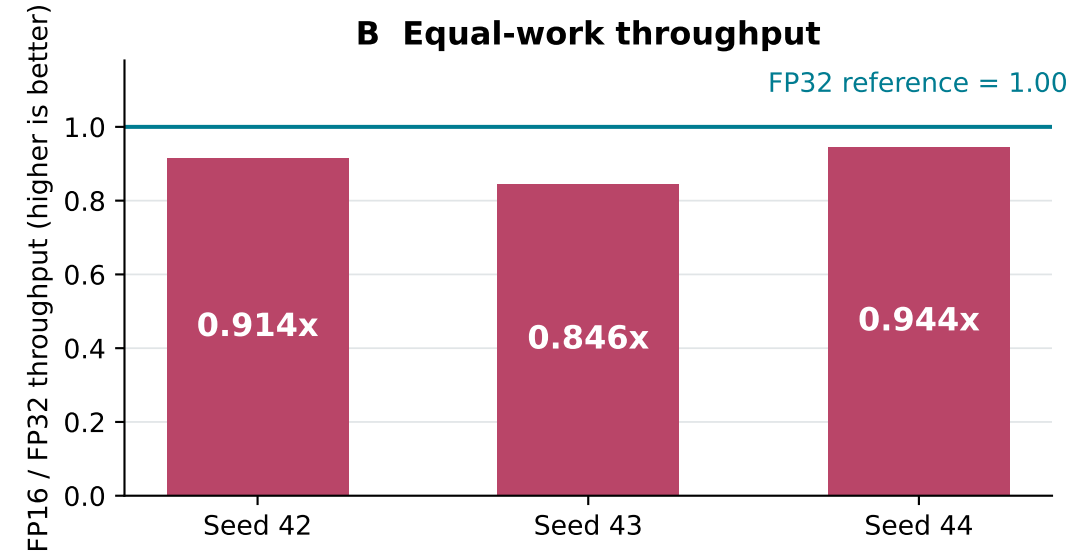
Frozen DeepSeek-produced model | 3 paired seeds | identical split and initial weights within pairs

Conservative: FP32

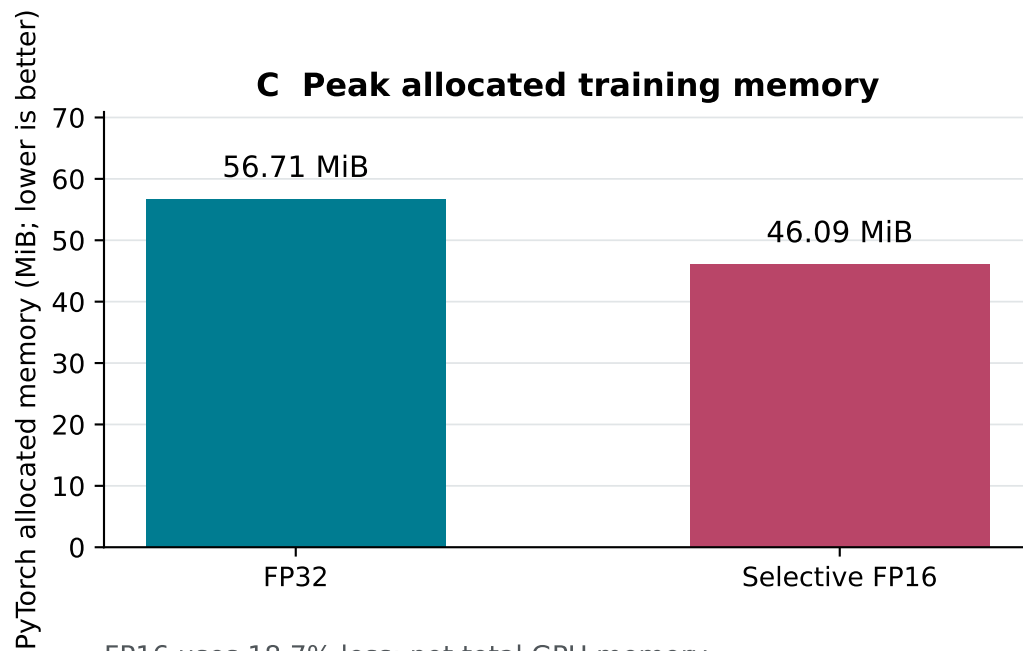
Normal: selective FP16



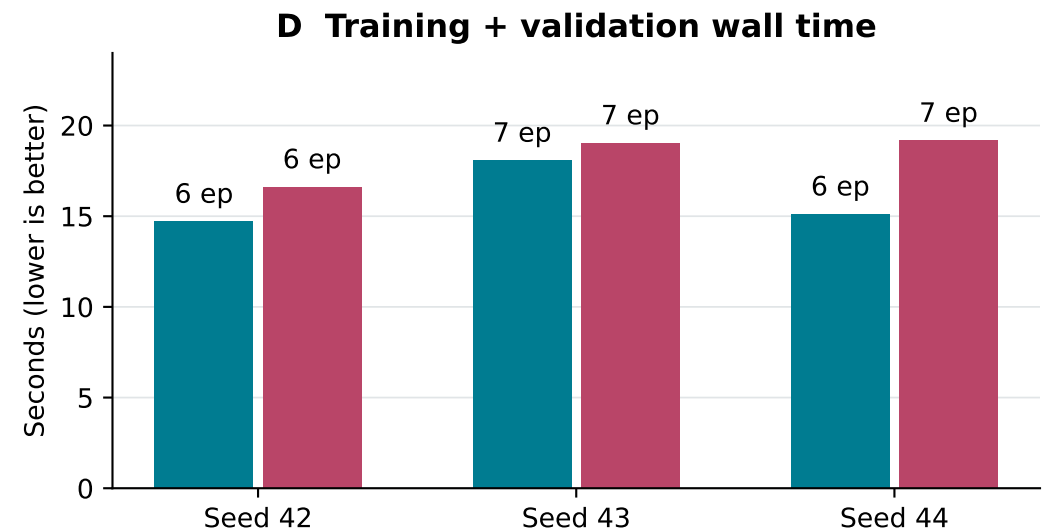
Zoomed F1 axis; labels identify initialization seeds.



Median: 8.6% lower with FP16; 100 timed updates.



FP16 uses 18.7% less; not total GPU memory.



Different stopping epochs; not an equal-work speed test.

Pilot only: one fixed internal split, not a leaderboard score or a general A10 recommendation.
Normal mode permits FP32; this experiment explicitly exercises its optional FP16 path.