

Reliability: what improved, and what remains unmeasured

Targeted regression fixes and runtime completion are different evidence from agent bug rate.

A Replayed against original source

	Original 93371dd	Updated 032609a
Recognize typed batch-size list	FAIL	PASS
Recognize typed learning-rate policy	FAIL	PASS
Preserve GPU flag on partial log update	FAIL	PASS

B Current frozen-model runs

	FP32	Selective FP16
	42	42
	43	43
	44	44

Selected checks failing: 3/3 before -> 0/3 now

Two defect areas; deliberately selected tests, not an unbiased bug sample.

Runtime failures: 0/6

All submissions verified; zero skipped updates.
One selected model, not six new agent designs.

Overall agent bug-rate reduction: NOT MEASURED

No matched original-versus-updated agent search exists for this experiment.
Earlier review rejections cannot be compared with these selected training runs as a failure-rate decrease.

Baseline: captured JustinLinKK hardware-awared checkout, not upstream main or latest remote HEAD.